



# ArceOS中 Mimalloc多线程实现 最终报告

邵志航

2024





## 从中期报告之后的四周做了什么样的工作？

- Week 5: 利用已有接口接入了一个新的测例并运行。
- Week 6-8: 阅读ArceOS与mimalloc源码，并积极尝试加入多线程支持。

先说结论：

就目前为止，我们仍找不到一个好的方法实现无锁多线程支持。  
但是也许有思路，只是仍需要时间去打磨和实现。

## Week 5: 利用已有接口接入了一个新的测例并运行。

```
/// the start function of the glibc_bench_test
pub fn glibc_bench_test_start(
    cb1: CallbackMalloc,
    cb2: CallbackMallocAligned,
    cb3: CallbackFree,
);
/// the start function of the glibc_bench_test_simple
pub fn glibc_bench_simple_test_start(
    cb1: CallbackMalloc,
    cb2: CallbackMallocAligned,
    cb3: CallbackFree,
);
/// the start function of the multi_thread_c_test
pub fn multi_thread_c_test_start(
    cb1: CallbackMalloc,
    cb2: CallbackMallocAligned,
    cb3: CallbackFree,
);
```

```
pub fn glibc_bench_simple_test() {
    println!("Glibc bench simple test begin...");
    let t0 = std::time::Instant::now();
    allocator_test::glibc_bench_simple_test();
    let t1 = std::time::Instant::now();
    println!("time: {:#?}", t1 - t0);
    println!("Glibc bench simple test OK!");
}
```

Week 6-8:


阅读ArceOS与mimalloc源码，并积极尝试加入多线程支持。

星期四 22:37

学长晚上好，就mimalloc多线程实现的事情又得麻烦你，我在global\_allocator里看到了mi\_alloc字段是Mutex<MiAllocator>，这是否意味着学长当时尚未完成的多线程实现是用这个Mutex先顶上，当多线程实现后达到的效果应该是去掉Mutex也可以正常运行内存分配器？

```
pub struct GlobalAllocator {  
    basic_alloc: Mutex<BasicAllocator<0>>,  
    buddy_alloc: Mutex<BuddyByteAllocator>,  
    slab_alloc: Mutex<SlabByteAllocator>,  
    tlsf_c_alloc: Mutex<TLSFCAllocator>,  
    tlsf_rust_alloc: Mutex<TLSFAllocator>,  
    mi_alloc: Mutex<MiAllocator>,  
    alloc_type: AllocType,  
    heap_address: usize,  
    heap_size: usize,  
}
```

应该是



```
pub struct GlobalAllocator {
    basic_alloc: Mutex<BasicAllocator<0>>,
    buddy_alloc: Mutex<BuddyByteAllocator>,
    slab_alloc: Mutex<SlabByteAllocator>,
    tlsf_c_alloc: Mutex<TLSFCAllocator>,
    tlsf_rust_alloc: Mutex<TLSFAllocator>,
    mi_alloc: MiAllocator,
    //mi_alloc: Mutex<MiAllocator>,
    alloc_type: AllocType,
    heap_arddress: usize,
    heap_size: usize,
}
```

```
running 1 test
mi alloc test:
Basic alloc test begin...
time: 246.45466ms
Basic alloc test OK!
Mi alloc test begin...
time: 41.782906ms
Mi alloc test OK!
Align alloc test begin...
time: 10.252158ms
Align alloc test OK!
Malloc large test begin...
time: 8.466428ms
Malloc large test OK!
Glibc bench test begin...
time: 4.770976691s
Glibc bench test OK!
Glibc bench simple test begin...
time: 12.394596604s
Glibc bench simple test OK!
Multi thread memory allocation test begin.
panicked at crates/allocator/tests/global_allocator.rs:161:48:
thread panicked while processing panic. aborting.
error: test failed, to rerun pass `--p allocator --test test_main`

Caused by:
  process didn't exit successfully: `./kerneldebug/arceos/target/release/deps/test_main-577ffbd2af5da6c --nocapture` (signal: 6, SIGABRT: process abort signal)
root@ab148eb3a638:/kerneldebug/arceos#
```

```

struct mi_heap_s {
    mi_tld_t*      tld;
    _Atomic(mi_block_t*) thread_delayed_free;
    mi_threadid_t  thread_id;           // thread this heap belongs too
    mi_arena_id_t  arena_id;           // arena id if the heap belongs to a specific arena (or 0)
    mi_random_ctx_t random;            // random number context used for secure allocation
    size_t         page_count;         // total number of pages in the `pages` queues.
    size_t         page_retired_min;    // smallest retired index (retired pages are fully free, but still
in the page queues)
    size_t         page_retired_max;    // largest retired index into the `pages` array.
    mi_heap_t*     next;               // list of heaps per thread
    bool           no_reclaim;         // `true` if this heap should not reclaim abandoned pages
    uint8_t        tag;               // custom tag, can be used for separating heaps based on the object
types
    mi_page_t*     pages_free_direct[MI_PAGES_DIRECT]; // optimize: array where every entry points a page with possibly
free blocks in the corresponding queue for that size.
    mi_page_queue_t pages[MI_BIN_FULL + 1];           // queue of pages for each size class (or "bin")
};

```

堆块具体结构的定义，其中堆块含有多个页，这符合先前实现里的定义。

```

/// mimalloc的heap结构
pub struct MiHeap {
    // page链表
    pub pages: [PagePointer; TOT_QUEUE_NUM],
    // thread id
    pub thread_id: usize,
    // thread delayed free
    pub thread_delayed_free: AtomicBlockPtr,
}

```

```

typedef struct mi_page_s {
    // "owned" by the segment
    uint8_t      segment_idx;    // index in the segment `pages` array, `page == &segment->pages[page->segment_idx]`
    uint8_t      segment_in_use:1; // `true` if the segment allocated this page
    uint8_t      is_committed:1;  // `true` if the page virtual memory is committed
    uint8_t      is_zero_init:1;  // `true` if the page was initially zero initialized
    uint8_t      is_huge:1;       // `true` if the page is in a huge segment

    // layout like this to optimize access in `mi_malloc` and `mi_free`
    uint16_t     capacity;        // number of blocks committed, must be the first field, see `segment.c:page_clear`
    uint16_t     reserved;        // number of blocks reserved in memory
    mi_page_flags_t flags;        // `in_full` and `has_aligned` flags (8 bits)
    uint8_t      free_is_zero:1;  // `true` if the blocks in the free list are zero initialized
    uint8_t      retire_expire:7; // expiration count for retired blocks

    mi_block_t*  free;            // list of available free blocks (`malloc` allocates from this list)
    mi_block_t*  local_free;      // list of deferred free blocks by this thread (migrates to `free`)
    uint16_t     used;            // number of blocks in use (including blocks in `thread_free`)
    uint8_t      block_size_shift; // if not zero, then `(1 << block_size_shift) == block_size` (only used for fast path)
} in `free.c:_mi_page_ptr_unalign`)
    uint8_t      heap_tag;        // tag of the owning heap, used to separate heaps by object type
    // padding

    size_t       block_size;      // size available in each block (always `>0`)
    uint8_t*     page_start;      // start of the page area containing the blocks

#ifdef (MI_ENCODE_FREELIST || MI_PADDING)
    uintptr_t    keys[2];         // two random keys to encode the free lists (see `_mi_block_next`) or padding canary
#endif

    _Atomic(mi_thread_free_t) xthread_free; // list of deferred free blocks freed by other threads
    _Atomic(uintptr_t)        xheap;

```

```

/// mimalloc的一个page控制头
#[derive(Clone)]
pub struct Page {
    // 块大小
    pub block_size: usize,
    // free链表
    pub free_list: BlockPointer,
    // local free list
    pub local_free_list: BlockPointer,
    // thread free list
    pub thread_free_list: AtomicBlockPtr,
    // page开始地址
    pub begin_addr: usize,
    // page结束地址
    pub end_addr: usize,
    // 尚未分配过的地址起点
    pub capacity: usize,
    // page链表中的上一项
    pub prev_page: PagePointer,
    // page链表中的下一项
    pub next_page: PagePointer,
    // 剩余块数
    pub free_blocks_num: usize,
}

```

```

typedef struct mi_segment_s {
    // constant fields
    mi_memid_t      memid;           // memory id to track provenance
    bool            allow_decommit;
    bool            allow_purge;
    size_t          segment_size;    // for huge pages this may be different from `MI_SEGMENT_SIZE`
    mi_subproc_t*   subproc;         // segment belongs to sub process

    // segment fields
    struct mi_segment_s* next;        // must be the first (non-constant) segment field -- see `segment.c:segment_init`
    struct mi_segment_s* prev;
    bool            was_reclaimed;    // true if it was reclaimed (used to limit on-free reclamation)

    size_t          abandoned;        // abandoned pages (i.e. the original owning thread stopped) (`abandoned <= used`)
    size_t          abandoned_visits; // count how often this segment is visited for reclaiming (to force reclaim if it is
too long)

    size_t          used;              // count of pages in use (`used <= capacity`)
    size_t          capacity;          // count of available pages (`#free + used`)
    size_t          segment_info_size; // space we are using from the first page for segment meta-data and possible guard
pages.
    uintptr_t       cookie;           // verify addresses in secure mode: `_mi_ptr_cookie(segment) == segment->cookie`

    struct mi_segment_s* abandoned_os_next; // only used for abandoned segments outside arena's, and only if
`mi_option_visit_abandoned` is enabled
    struct mi_segment_s* abandoned_os_prev;

    // layout like this to optimize access in `mi_free`
    _Atomic(mi_threadid_t) thread_id; // unique id of the thread owning this segment
    size_t                page_shift; // `1 << page_shift` == the page sizes == `page->block_size * page->reserved` (unless
the first page, then `-segment_info_size`).
    mi_page_kind_t        page_kind;   // kind of pages: small, medium, large, or huge
    mi_page_t             pages[1];    // up to `MI_SMALL_PAGES_PER_SEGMENT` pages
} mi_segment_t;


```

```

pub struct Segment {
    // 把mi_heap藏在第一个段的开头
    pub mi_heap: MiHeap,
    // page种类
    pub page_kind: PageKind,
    // 段的大小
    pub size: usize,
    // 包含多少个page
    pub num_pages: usize,
    // 每个page的头结构
    pub pages: [Page; MAX_PAGE_PER_SEGMEGT],
    /// thread_id
    pub thread_id: usize,
    // padding, 使空间对齐到8192
    pub padding: [usize; 431],
    // 接下来就是每个page的实际空间, 注意第一个page会小一些
}

```





```
running 1 test
mi alloc test:
Basic alloc test begin...
time: 2.829709758s
Basic alloc test OK!
Mi alloc test begin...
time: 10.936271344s
Mi alloc test OK!
Align alloc test begin...
thread 'test_start' panicked at crates/allocator_test/src/align_test.rs:33:13:
align not correct.
stack backtrace:
 0: rust_begin_unwind
   at /rustc/ed04567ba1d5956d1080fb8121caa005ce059e12/library/std/src/panicking.rs:665:5
 1: core::panicking::panic_fmt
   at /rustc/ed04567ba1d5956d1080fb8121caa005ce059e12/library/core/src/panicking.rs:74:14
 2: allocator_test::align_test::align_test
   at /kerneldebug/arceos/crates/allocator_test/src/align_test.rs:33:13
 3: test_main::align_test
   at ./tests/test_main.rs:83:5
 4: test_main::mi_alloc_test
   at ./tests/test_main.rs:395:5
 5: test_main::test_start
   at ./tests/test_main.rs:418:5
 6: test_main::test_start::{{closure}}
   at ./tests/test_main.rs:409:16
 7: core::ops::function::FnOnce::call_once
   at /rustc/ed04567ba1d5956d1080fb8121caa005ce059e12/library/core/src/ops/function.rs:250:5
 8: core::ops::function::FnOnce::call_once
   at /rustc/ed04567ba1d5956d1080fb8121caa005ce059e12/library/core/src/ops/function.rs:250:5
note: Some details are omitted, run with `RUST_BACKTRACE=full` for a verbose backtrace.
error: test failed, to rerun pass `-p allocator --test test_main`

Caused by:
  process didn't exit successfully: `/kerneldebug/arceos/target/debug/deps/test_main-046ad87eb4743645 --nocapture` (signal: 11, SIGSEGV: invalid memory reference)
```

```


error[E0204]: the trait `core::marker::Copy` cannot be implemented for this type
--> crates/mimalloc_allocator/src/data.rs:62:17
62 | #[derive(Clone, Copy)]
   |               ^^^^
...
71 |     pub thread_free_list: AtomicPtr<BlockPointer>,
   |     ----- this field does not implement `core::marker::Copy`
= note: this error originates in the derive macro `Copy` (in Nightly builds, run with -Z macro-backtrace for more info)

error[E0277]: the trait bound `AtomicPtr<data::BlockPointer>: Clone` is not satisfied
--> crates/mimalloc_allocator/src/data.rs:71:5
62 | #[derive(Clone, Copy)]
   |         ----- in this derive macro expansion
...
71 |     pub thread_free_list: AtomicPtr<BlockPointer>,
   |     ~~~~~~ the trait `Clone` is not implemented for `AtomicPtr<data::BlockPointer>`

note: required by a bound in `AssertParamIsClone`
--> /root/.rustup/toolchains/nightly-x86_64-unknown-linux-gnu/lib/rustlib/src/rust/library/core/src/clone.rs:198:34
198 | pub struct AssertParamIsClone<T: Clone + ?Sized> {
   |         ~~~~~~ required by this bound in `AssertParamIsClone`
= note: this error originates in the derive macro `Clone` (in Nightly builds, run with -Z macro-backtrace for more info)

Some errors have detailed explanations: E0204, E0277.
For more information about an error, try `rustc --explain E0204`.
error: could not compile `mimalloc_allocator` (lib) due to 2 previous errors

```




```
pub struct AtomicBlockPtr {
    pub ptr: AtomicPtr<BlockPointer>,
}

impl AtomicBlockPtr {
    fn new(ptr: *mut BlockPointer) -> Self {
        AtomicBlockPtr {
            ptr: AtomicPtr::new(ptr)
        }
    }
}

impl Clone for AtomicBlockPtr {
    fn clone(&self) -> Self {
        let atomic_ptr_clone = self.ptr.load(Ordering::SeqCst);
        let new_atomic_ptr = AtomicPtr::new(atomic_ptr_clone);

        AtomicBlockPtr {
            ptr: AtomicPtr::new(atomic_ptr_clone)
        }
    }
}
```



```

error[E0596]: cannot borrow `self.mi_alloc` as mutable, as it is behind a `&` reference
--> crates/allocator/tests/global_allocator.rs:164:34
164 |         if let Ok(ptr) = self.mi_alloc.inner_mut().alloc(size, align_pow2) {
            ^^^^^^^^^^^^^ `self` is a `&` reference, so the data it refers to cannot be borrowed as mutable
help: consider changing this to be a mutable reference
119 |         pub unsafe fn alloc(&mut self, layout: Layout) -> AllocResult<usize> {
            ^^^^^^^^^^^^^
error[E0596]: cannot borrow `self.mi_alloc` as mutable, as it is behind a `&` reference
--> crates/allocator/tests/global_allocator.rs:213:17
213 |         self.mi_alloc.inner_mut().dealloc(pos, size, align_pow2);
            ^^^^^^^^^^^^^ `self` is a `&` reference, so the data it refers to cannot be borrowed as mutable
help: consider changing this to be a mutable reference
174 |         pub unsafe fn dealloc(&mut self, pos: usize, layout: Layout) {
            ^^^^^^^^^^^^^
For more information about this error, try `rustc --explain E0596`.
error: could not compile `allocator` (test "test_main") due to 2 previous errors

```



几个解决方案：

1. `RefCell<T>`：提供内部可变性。

但是这样的方法行不通。因为会有 `Already borrowed: BorrowMutError`。这是很正常的事情，因为 `RefCell` 本身就不保证线程安全。

2. `Miallocator` 的多线程改造。

需要参考 `RefCell` 进行操作，但是又不能引入锁。



```
type BorrowFlag = AtomicIsize;
/// A byte-granularity memory allocator based on the [mimalloc_allocator] written by rust
pub struct MiAllocator {
    //Atomic...
    borrow: BorrowFlag,
    data: UnsafeCell<MiAllocatorInner>,
}

pub struct MiAllocatorInner {
    inner: Option<Heap>,
}

impl MiAllocator {
    /// Creates a new empty `TLSFAllocator`.
    pub const fn new() -> Self {
        Self {
            borrow: AtomicIsize::new(0),
            data: UnsafeCell::new(MiAllocatorInner::new())
        }
    }

    pub fn inner_mut(&mut self) -> &mut MiAllocatorInner {
        self.data.get_mut()
    }

    pub fn inner(&self) -> &MiAllocatorInner {
        let ptr = self.data.get();
        let reference: &MiAllocatorInner = unsafe {
            &*ptr
        };
        reference
    }
}
```



未来的工作：

- 1.至少先完善数据结构，能够运行起多线程测例。
- 2.现有增加的mimalloc内部字段如何使用仍需要探索。

有鉴于此，也许这需要另外的八周去完成，也许需要更长的时间。

The background features abstract geometric shapes. On the left, there is a large grey parallelogram and a green parallelogram, both tilted. Several thin green lines are scattered around these shapes. On the right, there is a large green parallelogram, also tilted, with a thin green line extending from its right side.

# Thanks

---