

对Mimalloc 在ArceOS下 线程安全实现的一些思考


邵志航

2024



思考：Mimalloc如何实现线程安全

- Thread-local heaps
- Lock-free data structures
- Deferred reclamation
- Batch Processing




“mimalloc 中的页归属于线程局部的堆，其分配是本地完成的，因此无锁是理所当然的。但是，对于一个对象来说，任意线程都可以释放它。为了避免 free() 操作持锁，mimalloc 再一次为每页面切分 thread free list，并且把其他线程释放的对象都使用原子操作放到这个链表中。”

ini

 代码解读 复制代码

```
1 void free( void* p ) {
2     //找到p所在的segment
3     segment_t* segment = (segment_t*)((uintptr_t)p & ~(4*MB));
4     if (segment==NULL) return;
5     //找到p所在的page
6     page_t* page = &segment->pages[(p - segment) >> segment->page_shift];
7     block_t* block = (block_t*)p;
8
9     if (thread_id() == segment->thread_id) { // free的是本线程分配的内存 (local free)
10         block->next = page->local_free;
11         page->local_free = block;
12         page->used--;
13         if (page->used - page->thread_freed == 0) page_free(page);
14     }
15     else { // free的是其他线程分配的内存 (non-local free)
16         atomic_push( &page->thread_free, block);
17         atomic_incr( &page->thread_freed );
18     }
19 }
```



```
// Free a block
// Fast path written carefully to prevent register spilling on the stack
void mi_free(void* p) mi_attr_noexcept
{
    mi_segment_t* const segment = mi_checked_ptr_segment(p, "mi_free");
    if mi_unlikely(segment==NULL) return;

    const bool is_local = (_mi_prim_thread_id() == mi_atomic_load_relaxed(&segment->thread_id));
    mi_page_t* const page = _mi_segment_page_of(segment, p);

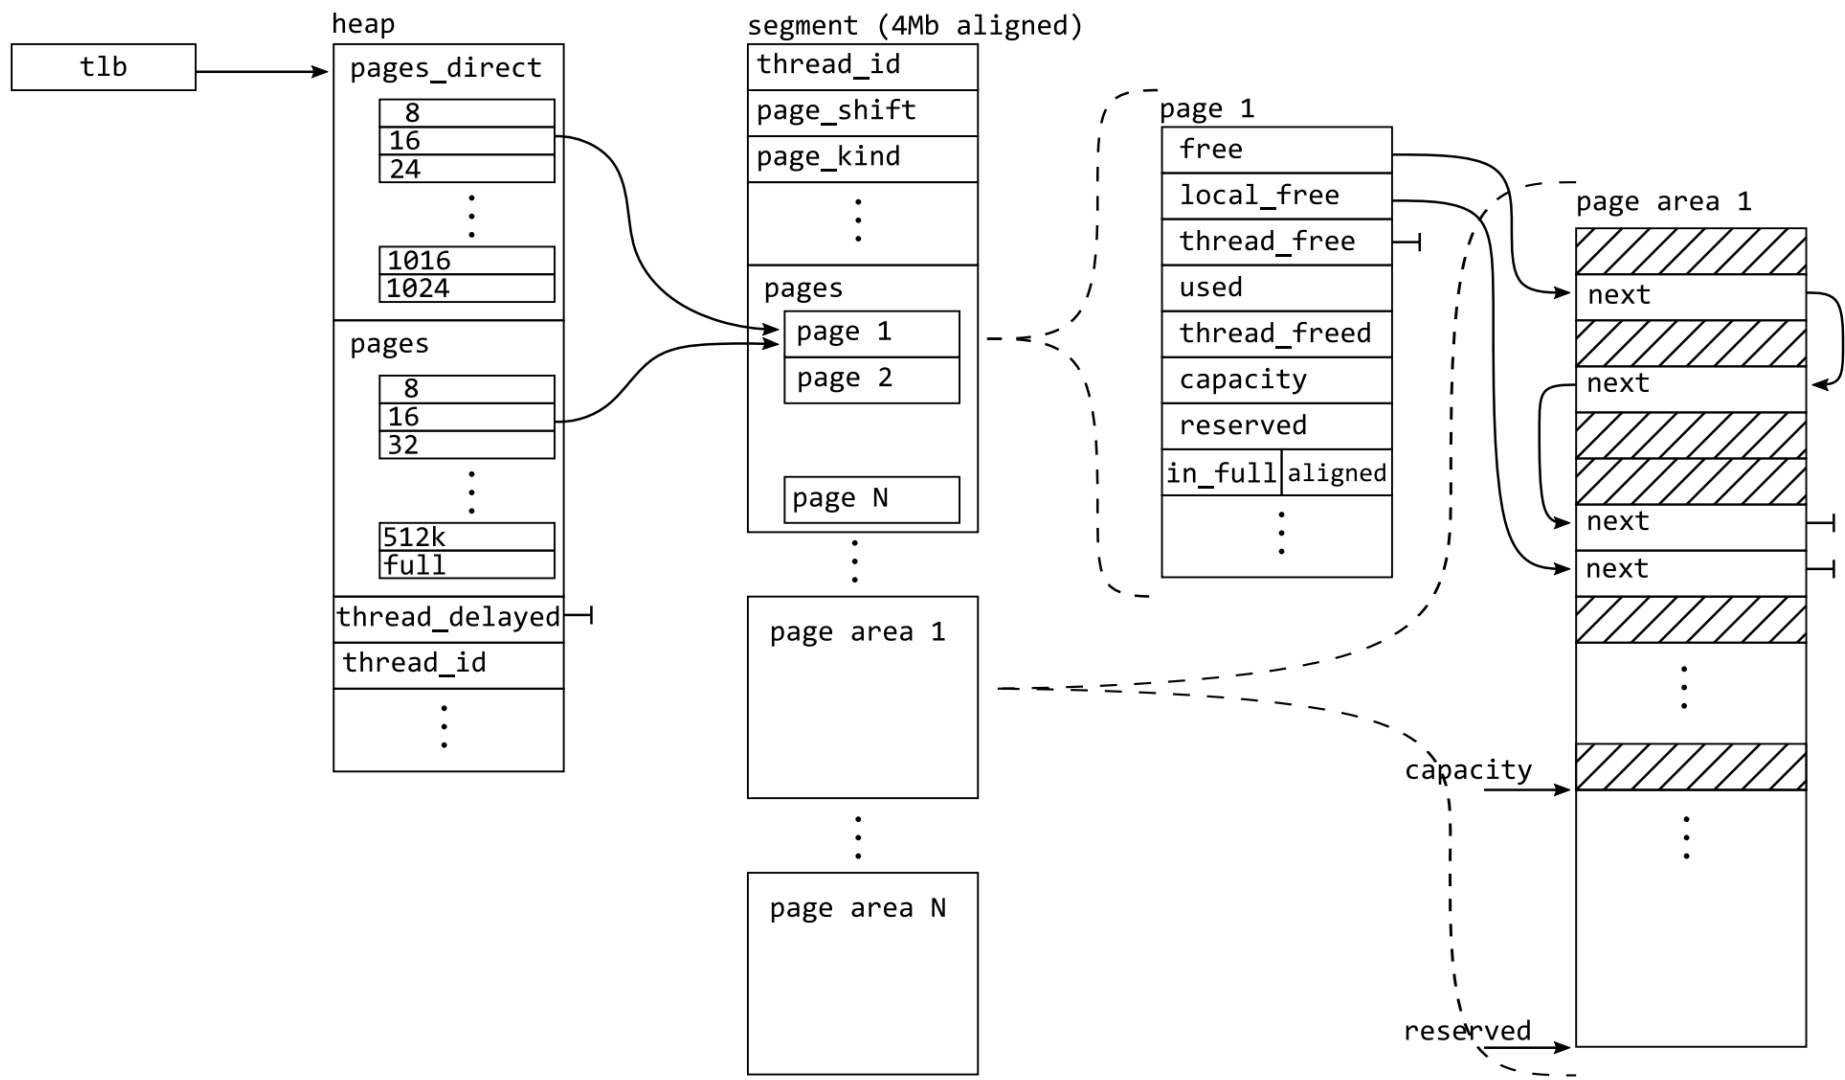
    if mi_likely(is_local) {
        // thread-local free?
        if mi_likely(page->flags.full_aligned == 0) { // and it is not a full page (full pages need to move from the full bin), nor has aligned blocks (aligned blocks need to be unaligned)
            // thread-local, aligned, and not a full page
            mi_block_t* const block = (mi_block_t*)p;
            mi_free_block_local(page, block, true /* track stats */, false /* no need to check if the page is full */);
        }
        else {
            // page is full or contains (inner) aligned blocks; use generic path
            mi_free_generic_local(page, segment, p);
        }
    }
    else {
        // not thread-local; use generic path
        mi_free_generic_mt(page, segment, p);
    }
}
```


ArceOS下已有的工作：

```
pub struct GlobalAllocator {
    basic_alloc: Mutex<BasicAllocator<0>>,
    buddy_alloc: Mutex<BuddyByteAllocator>,
    slab_alloc: Mutex<SlabByteAllocator>,
    tlsf_c_alloc: Mutex<TLSFCAllocator>,
    tlsf_rust_alloc: Mutex<TLSFAllocator>,
    mi_alloc: Mutex<MiAllocator>,
    alloc_type: AllocType,
    heap_address: usize,
    heap_size: usize,
}

const PAGE_SIZE: usize = 1 << 22; // need 4MB aligned (prepared for mimalloc)
const HEAP_SIZE: usize = 1 << 26; // 512MB
/// The memory heap in this test
static mut HEAP: [usize; HEAP_SIZE + PAGE_SIZE] = [0; HEAP_SIZE + PAGE_SIZE];

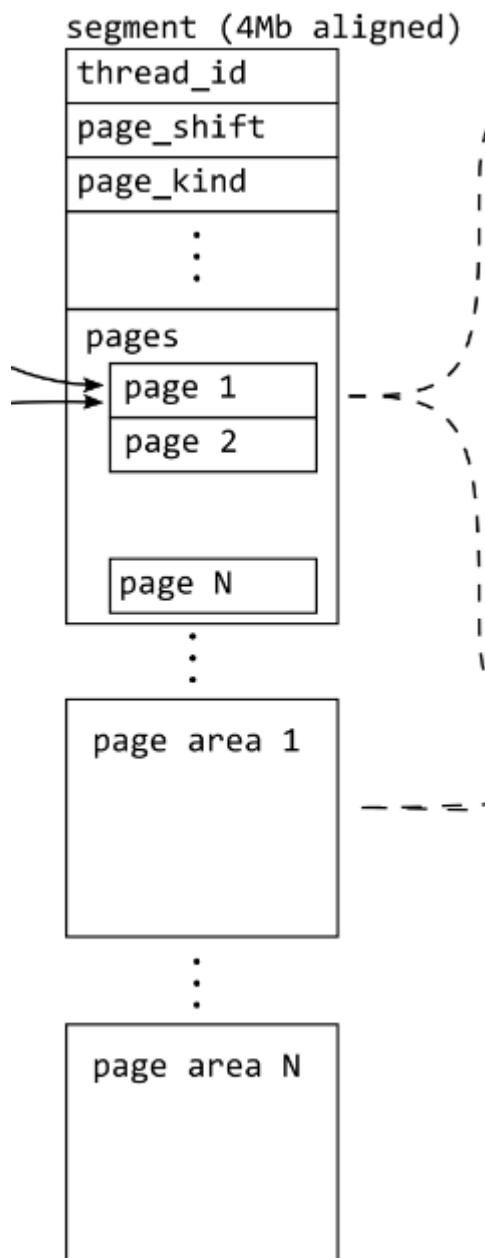
/// The global allocator to test alloc in user mode.
/// The alloc mode supported: system_alloc, basic_alloc, buddy_alloc,
/// slab_alloc, tlsf_c_alloc and tlsf_rust_alloc.
/// The basic_alloc can choose strategy: first_fit, best_fit and worst_fit.
impl GlobalAllocator {
    pub const fn new() -> Self {
        Self {
            basic_alloc: Mutex::new(BasicAllocator::new()),
            buddy_alloc: Mutex::new(BuddyByteAllocator::new()),
            slab_alloc: Mutex::new(SlabByteAllocator::new()),
            tlsf_c_alloc: Mutex::new(TLSFCAllocator::new()),
            tlsf_rust_alloc: Mutex::new(TLSFAllocator::new()),
            mi_alloc: Mutex::new(MiAllocator::new()),
            alloc_type: AllocType::SystemAlloc,
            heap_address: 0,
            heap_size: 0,
        }
    }
}
```






/// mimalloc的heap结构

```
pub struct MiHeap {  
    // page链表  
    pub pages: [PagePointer; TOT_QUEUE_NUM],  
}
```

```
pub struct Segment {  
    // 把mi_heap藏在第一个段的开头  
    pub mi_heap: MiHeap,  
    // page种类  
    pub page_kind: PageKind,  
    // 段的大小  
    pub size: usize,  
    // 包含多少个page  
    pub num_pages: usize,  
    // 每个page的头结构  
    pub pages: [Page; MAX_PAGE_PER_SEGMEGT],  
    // padding, 使空间对齐到8192  
    pub padding: [usize; 434],  
    // 接下来就是每个page的实际空间, 注意第一个page会小一些  
}
```



```
/// mimalloc的一个page控制头
#[derive(Clone, Copy)]
pub struct Page {
    // 块大小
    pub block_size: usize,
    // free链表
    pub free_list: BlockPointer,
    // page开始地址
    pub begin_addr: usize,
    // page结束地址
    pub end_addr: usize,
    // 尚未分配过的地址起点
    pub capacity: usize,
    // page链表中的上一项
    pub prev_page: PagePointer,
    // page链表中的下一项
    pub next_page: PagePointer,
    // 剩余块数
    pub free_blocks_num: usize,
}
```



TODO:

- ▶ 将C语言实现的Mimalloc官方仓库接入arceos
 - ▶ <https://github.com/microsoft/mimalloc>
- ▶ 实现Rust语言、支持多线程的Mimalloc
- ▶ 增加更多测试多线程的测例

接入C语言实现的Mimalloc库，可参考：

```
fn main() {  
    let out_dir: String = env::var(key: "OUT_DIR").unwrap();  
  
    Command::new(program: "cc").args(&["tests/test.c", "-O3", "-c", "-fPIC", "-o"]) &mut Command  
        .arg(&format!("{}/test.o", out_dir)) &mut Command  
        .status().unwrap();  
  
    Command::new(program: "ar").args(&["crus", "libtest.a", "test.o"]) &mut Command  
        .current_dir(&Path::new(&out_dir)) &mut Command  
        .status().unwrap();  
  
    println!("cargo:rustc-link-search=native={}", out_dir);  
    println!("cargo:rustc-link-lib=static=test");  
}
```

```
use std::ffi::c_ulonglong;  
#[link(name = "tlsf")]  
extern {  
    pub fn tlsf_create_with_pool(mem: c_ulonglong, bytes: c_ulonglong) -> c_ulonglong;  
    pub fn tlsf_add_pool(tlsf: c_ulonglong, mem: c_ulonglong, bytes: c_ulonglong) -> c_ulonglong;  
  
    pub fn tlsf_malloc(tlsf: c_ulonglong, bytes: c_ulonglong) -> c_ulonglong; //申请一段内存  
    pub fn tlsf_memalign(tlsf: c_ulonglong, align: c_ulonglong, bytes: c_ulonglong) -> c_ulonglong; //申请一段内存，要求对齐到align  
    pub fn tlsf_free(tlsf: c_ulonglong, ptr: c_ulonglong); //回收  
}
```



增加更多测试多线程的用例，计划移植：

<https://github.com/daanx/mimalloc-bench/blob/master/bench/glibc-bench/bench-malloc-thread.c>

References:

1. <https://github.com/daanx/mimalloc-bench>
2. <https://github.com/microsoft/mimalloc>
3. <https://www.bluepuni.com/archives/paper-reading-mimalloc-free-list-sharding-in-action>
4. [mimalloc-tr-v1.pdf](#)
5. <https://www.cnblogs.com/Five100Miles/p/12169392.html>

The background features abstract geometric shapes. On the left, there is a large grey parallelogram and a green parallelogram, both tilted. Several thin green lines are scattered around these shapes. On the right, there is a large green parallelogram, also tilted, with a thin green line nearby. The word "Thanks" is centered in the middle of the slide.

Thanks
